


Where the alert sits, and what happens after someone reads it

The analysis is built and running. This is the operating workflow around it — who receives a finding, what they do with it, what the system does in response, and how the loop closes.

1 · The operating loop

Solid boxes exist and run today. Dashed boxes are the workflow around them — the part that turns an alert into an action with an owner and an outcome.



The loop today stops at the Teams card. Everything dashed — the change trigger, what an Approve produces, the answer path, and the verification that closes the loop — is the workflow that has to exist for the alert to sit inside a solution rather than beside one.

— BUILT AND RUNNING - - - TO BUILD

THE GAP, STATED PLAINLY

Today a finding is *read*. Nothing records that it was read, nobody is assigned, and next Monday the same finding is posted again with a slightly larger number — because unfulfilled tickets keep accruing. Without the dashed half, the alert repeats rather than resolves.

2 • What Approve and Reject actually do

This is the decision that unblocks everything else, and it is not ours to make alone. Approve cannot mean "raise the threshold automatically" — that is a capacity policy change requiring authority the agent does not have, and the design explicitly forbids auto-execution.

OPTION	ON APPROVE	NEEDS	EFFORT
A • Acknowledge	Log who approved and when; @mention the region owner in-thread; suppress the repeat for N cycles	Nothing beyond Teams	Small
B • Raise a work item (recommended)	Create an ADO / ICM item with the region, the sized threshold, the evidence incidents and a due date; link it back into the thread	ADO or ICM project + service account	Medium
C • Draft the change	Generate the actual capacity-policy change request, routed for a second approval before submission	Access to the capacity-request system	Large

B is what "the capacity team who's working on it" implies — Approve should put something in someone's queue, with the sizing and evidence already attached, so the work does not have to be re-derived from the message.

Reject is not a dead end

A rejection carries a reason, and the reason is a signal worth keeping:

- "Threshold is already being raised" → suppress this region until the next quarter
- "That customer is churning" → exclude the subscription from exposure
- "48h is the wrong cut-off for Standard tier" → change the SLA in config, which re-scores every ticket

Recorded rejections are how the alert stops crying wolf. Without them it repeats the same finding indefinitely.

3 • The card, restructured

The current card states the problem, cause, impact and action — but in prose, so the number has to be read for. Same content, fixed fields, scannable in five seconds.

\$146K of customer revenue is exposed to capacity denials. 30 of 60 requests failed this period. Three regions carry 88% of it.

For: Capacity Operations · FYI: Finance & Sales leadership

westeurope**\$76,309**

PROBLEM	4 of 9 requests failed — 3 approved late (median 10 days), 1 never fulfilled
CAUSE	Approval latency, not capacity shortage. Every request was eventually granted
IMPACT	3 customers · \$690K ARR affected · \$76,309 at risk
ACTION	Raise the auto-approval ceiling to 166 units
EFFECT	Would have cleared 100% of these without review
EVIDENCE	INC 677681988 — 83u, 25 days late, \$37K
OWNER	Capacity Operations — westeurope

northcentralus \$42,352 · eastus2 \$10,785 — same structure, two more blocks.
Figures illustrative: ARR reference is placeholder. Rankings hold; dollars do not.

Approve

Reject

Ask

Every field is already computed — this is a rendering change, not new analysis. The one new element is **Owner**, which needs a region-to-team mapping we do not currently have.

4 · Keeping the alert true when things change

His question — *"if thresholds get increased or cores get increased, is my alert calling out the same thing?"* — is the sharpest one asked, and today the answer is no.

WHAT CHANGES	TODAY	SHOULD BE
New tickets arrive in ICM	Picked up next Monday	Trigger a run on arrival
A threshold is actually raised	Not detected — the same finding reappears	Detected; the finding closes and the loop reports the drop in exposure
The SLA cut-off is changed	Needs a config edit and a manual re-run	Config change re-scores every ticket automatically
An unfulfilled ticket is finally granted	Silently stops accruing	Reported as resolved, with the final cost of the delay

PER-TIER SLA – A CHANGE TO THE MODEL

He described **different cut-offs per subscription tier** — 48h for Enterprise, 24h or 32h for others. Today there is a single global 48-hour cut-off applied to everyone. Making it per-tier changes which tickets count as failures, and therefore every number in the report. This needs the real SLA table before it can be implemented.

5 • What we need decided

DECISION 1

What does Approve produce?

Acknowledge in-thread (A), raise a work item (B), or draft the policy change (C). If B, which system — ADO or ICM — and which project.

DECISION 2

What is the SLA per subscription tier?

Enterprise 48h was mentioned, with 24h and 32h for others. The real table changes what counts as a failure and every figure downstream.

DECISION 3

Who owns each region?

The card should name an owner. A region-to-team mapping does not exist yet — without it, "Owner" stays generic and nobody is actually accountable for a finding.

DECISION 4

Should the agent answer questions in the channel?

The engine is built — ranking, totals, per-customer, per-incident, methodology, all answered from the same computed tables. Making it two-way needs a registered Teams bot rather than the current one-way webhook. That is an integration and an app registration, not new analysis.